Lab 6

In this lab, you should perform **task 0 and task 1 before attending the lab session**.

> Lab 6 is built on top of the code infrastructure of Lab 5, i.e. the "shell". Naturally, you are expected to use the code you wrote for the previous lab.

Motivation

In this lab you will enrich the set of capabilities of your shell by implementing **input/output redirection** and **pipelines** (see the reading material). Your shell will then be able to execute non-trivial commands such as **"tail -n 2 in.txt| cat > out.txt"**, demonstrating the power of these simple concepts.

Lab 6 tasks

Through out the tasks of the lab, you are asked to add debug messages, according to the specifications if given, and as you see fit where no specification is provided.

Task 0

Pipes

A pipe is a pair of input stream/output stream, such that one stream feeds the other stream directly. All data that is written to one side (the "write end") can be read from the other side (the "read end"). This sort of feed becomes pretty useful when one wishes to communicate between processes.

Your task: Implement a simple program called **mypipe**, which creates a child process that sends the message "hello" to its parent process. The parent then prints the incoming message and terminates. Use the **pipe** system call (see man) to create the pipe.

Task 1

Redirection

Add standard input/output redirection capabilities to your shell (e.g. **"cat < in.txt > out.txt"**). Guidelines on I/O redirection can be found in the reading material.

Notes:

- The **inputRedirect** and **outputRedirect** fields in cmdLine do the parsing work for you. They hold the redirection file names if exist, NULL otherwise.
- Remember to redirect input/output only in the child process. We do not want to redirect the I/O of the shell itself (parent process).

Task 2

Here we wish to save internal variables in our shell. This is done by adding a list of string pairs (name, value), which associates names with values. When variable names appear in the command line with a "$" prefix, they are replaced with their associated value.
For instance, let variable **i** be mapped to value **Hello**. "`echo $i`" will then print "Hello", whereas "`echo i`" will print "i". Similarly, the command "`ls $i`" will be translated to "ls Hello", etc.

The task is divided into four mini-tasks, which define operations relevant to the internal variables. Although this can be implemented in several ways, **you are strictly required to base your implementation on a linked list of (name, value) string pairs**. Your implementation should be general enough to support unlimited number of variables and unlimited length of names and values.

1. Add a *set* command to the shell, which associates a given name with a given value. **If the name already exists, the command should override the existing value**.
   Usage: "set x y" creates a variable with name x and value y.

2. Add a *vars* command, which prints all variables in the internal variables list.
   Usage: "vars" prints out all variables.

3. Activate the internal variables in each executed command line, by replacing each argument that starts with a **$** sign with its proper value. **Write an appropriate error message when variables are not found**.
   Hint: Use the `int replaceCmdArg(cmdLine *pCmdLine, int num, const char *newString);` function from LineParser.c and LineParser.h from Lab 5

4. Adding support of ~ for `cd`. Initiating the command `cd` ~ in your shell will result in an `error`. You need to add support for converting ~ to your home directory. Find the variable for your home directly, **do not hard code it**.
   Hint: Refer to the **getenv(3)** for help

> **You will need to propagate the error messages to stderr, such as:**
> - Activating a variable that does not exist.

Task 3

Note
Task 3 is independent of the shell we revisited in task 1. You're not allowed to use the LineParser functions in this task. However, you need to declare an array of strings containing all of the arguments and ending with 0 to pass to execvp() just like the one returned by parseCmdLines().

Here we wish to explore the implementation of a pipeline. In order to achieve such a pipeline, one has to create pipes and properly redirect the standard outputs and standard inputs of the processes.
Please refer to the 'Introduction to Pipelines' section in the reading material.

Your task: Write a short program called **mypipeline** which creates a pipeline of 2 child processes. Essentially, you will implement the shell call **"ls -l | tail -n 2"**.
(A question: what does "ls -l" do, what does "tail -n 2" do, and what should their combination produce?)

Follow the given steps as closely as possible to avoid synchronization problems:

1. Create a pipe.
2. Fork to a child process (child1).
3. On the child1 process:

    a. Close the standard output.
    b. Duplicate the write-end of the pipe using **dup** (see man).
    c. Close the file descriptor that was duplicated.
    d. Execute "ls -l".

4. **On the parent process: Close the write end of the pipe.**
5. Fork again to a child process (child2).
6. On the child2 process:

    a. Close the standard input.
    b. Duplicate the read-end of the pipe using **dup**.
    c. Close the file descriptor that was duplicated.
    d. Execute "tail -n 2".

7. **On the parent process: Close the read end of the pipe.**
8. Now wait for the child processes to terminate, in the same order of their execution.

Mandatory Requirements

1. Compile and run the code and make sure it does what it's supposed to do.
2. Your program must print the following debugging messages if the argument -d is provided. All debugging messages must be sent to stderr! These are the messages that should be added:

- On the parent process:
  - Before forking, "(parent_process>forking…)"
  - After forking, "(parent_process>created process with id:)"
  - Before closing the write end of the pipe, "(parent_process>closing the write end of the pipe…)"
  - Before closing the read end of the pipe, "(parent_process>closing the read end of the pipe…)"
  - Before waiting for child processes to terminate, "(parent_process>waiting for child processes to terminate…)"
  - Before exiting, "(parent_process>exiting…)"
- On the 1st child process:
  - "(child1>redirecting stdout to the write end of the pipe…)"
  - "(child1>going to execute cmd: …)"
- On the 2nd child process:
  - "(child2>redirecting stdin to the read end of the pipe…)"
  - "(child2>going to execute cmd: …)"

3. How does the following affect your program:

   a. Comment out step 4 in your code (i.e. on the parent process:**do not** Close the write end of the pipe). Compile and run your code. (Also: see "man 7 pipe")
   b. Undo the change from the last step. Comment out step 7 in your code. Compile and run your code.
   c. Undo the change from the last step. Comment out step 4 and step 8 in your code. Compile and run your code.

Task 4

Go back to your shell and add support to a single pipe. Your shell must be able now to run commands like: `ls|wc -l` which basically counts the number of files/directories under the current working dir. The most important thing to remember about pipes is that the write-end of the pipe needs to be closed in all processes, otherwise the read-end of the pipe will not receive EOF, unless the main process terminates.

Deliverables:

Tasks 1,2,3 must be completed during the regular lab. Task 4 may be done in a completion lab, but only if you run out of time during the regular lab. The deliverables must be submitted until the end of the lab session.
You must submit source files for task 1, task 2, task 3, and task 4 and a makefile that compiles them. The source files must be named task1.c, task2.c, task3.c, task4.c, and makefile

Submission instructions

- Create a zip file with the relevant files (only).
- Upload zip file to the submission system.

- Download the zip file from the submission system and extract its content to an empty folder.
- Compile and test the code to make sure that it still works.

Example structure of zip, + is a directory and - is a file:

```
+ t1
  - makefile
  - task1.c
+ t2
  - makefile
  - task2.c
+ t3
  - makefile
  - task3.c
+ t4
  - makefile
  - task4.c
```