Lab 7: ELF-introduction

This lab **may** be done in pairs.

You must complete tasks 0, 1a, 1b, 1c before attending the lab.

Lab Goals

- Extracting useful information from files in ELF.
- Fixing files using this information: reverse engineering.

In the following labs, you will learn to handle object and executable files. We will begin by learning just some of the basics of ELF, together with applications you can already use at this level - editing binary files and writing software patches. Then, we will continue our study of ELF files, by beginning to parse the structures of ELF files, and to use them for various purposes. In particular, we will access the data in the section header table and in the symbol table.

Methodology

- [Get to know the ELF](#).
- Learn how to use the *readelf* utility. By using *readelf* you can get, in a human readable format, all the ELF structural information.
- Experience basic ELF manipulation.

Recommended Operating Procedure

This advice is relevant for all tasks. Note that while at some point you will no longer be using *hexedit* to process the file and *readelf* to get the information, nevertheless in some cases you may still want to use these tools for debugging purposes. In order to take advantage of these tools and make your tasks easier, you should:

- Support debugging messages: in particular the offsets of the various items, as you discover them from the headers. Also, whenever the user is required to enter values, you should print the **parsed** values in their respective representation (e.g. string, decimal or hexadecimal).
- Use *hexedit* and *readelf* to compare the information you are looking for, especially if you run into unknown problems: *hexedit* is great if you know the exact location of the item you are looking for.
- Note that while the object files you will be processing will be linked using *ld*, and will, in most cases, use direct system calls in order to make the ELF file simpler, there is no reason why the programs you write need use this interface. You are allowed to use the standard library when building your own C programs.

- In order to preserve your sanity, even if the code you MANIPULATE may be without stdlib, we advise that for your OWN CODE you DO use the C standard library! (Yes, this is repeated twice, so that you notice it!)
- In order to keep sane in the following labs as well, **understand** what you are doing and **keep track** of that and of your code, as you will be using them in future labs.

All the executable files we will work with in this session are 32-Bit ELF binaries.
Compile your code accordingly.

Lab 7 Tasks

Deliverables
You should read and understand the reading material, and do tasks 0, 1a, 1b, 1c before attending the lab. To be eligible for a full grade, you must complete all of tasks up-to (and including) task 3a, during the regular lab. Tasks 3b and 4 may be done in a completion lab, if you run out of time.

Task 0

Task 0a:

Download the following file: a.out. Answer the following questions (be prepared to explain your answers to the lab instructor):

1. Where is the entry point specified, and what is its value?
2. How many sections are there in a.out?
3. What is the size of the .text section?
4. Does the symbol _start occur in the file? If so, where is it mapped to in virtual memory?
5. Does the symbol main occur in the file? If so, where is it mapped to in virtual memory?
6. Where in the file does the code of function "main" start?

Task 0b

Write a program called *hexeditplus*:

```
./hexeditplus
```

The hexeditplus program performs operations (read and write) on files and memory. File operations are done on a file *file_name* as defined below. Each operation is done in units of *size* bytes, which indicates a unit size, i.e. the number of bytes we want to use as the basic unit in each operation of our program, such as "display memory contents". Size can be either 1, 2 or 4, with 1 as the default.

First, define a menu for the user with a number of predefined functions (as done in Lab 2), to which we will add functions as we go. The program prints the menu, obtains a choice from the user, acts on it, and repeats infinitely. For example, if the functions: "Toggle Debug Mode", "Set File Name", "Set Unit Size" and "Quit" are available, then the command line:

```
 ./hexeditplus
```

Will print:

```
Choose action:
0-Toggle Debug Mode
1-Set File Name
2-Set Unit Size
3-Quit
```

For this part, use an array with the above menu names and pointers to appropriate functions that implement each option.

At this point implement "Toggle Debug Mode", "Set File Name", "Set Unit Size", and "Quit".

Toggle debug mode means turn the debug flag on (if it is currently off, which it is in the initial state), and print "Debug flag now on". If the debug flag is on, this function prints "Debug flag now off", and turns the flag off. When the debug mode is on, you should print the value of the variables: unit_size, file_name, and mem_count, every time just before the menu is printed.

All functions should be of the form:

```
 void fun(state* s); // Getting the state as a pointer allows the functions to change it.
```

where the state struct is defined as:

```
typedef struct {
   char debug_mode;
   char file_name[128];
   int unit_size;
   unsigned char mem_buf[10000];
   size_t mem_count;
   /*
    .

    .
    Any additional fields you deem necessary
   */
} state;
```

Set File Name queries the user for a file name, and store it in *file_name*. You may assume that the file name is no longer than 100 characters. If debug mode is on, the function should also print: "Debug: file name set to 'file_name' " (obviously, replacing 'file_name' with the actual name).

Set Unit Size option sets the size variable. The steps are:

1. Prompt the user for a number.
2. If the value is valid (1, 2, or 4), set the size variable accordingly.
3. If debug mode is on, print "Debug: set size to x", with x the appropriate size.
4. If not valid, print an error message and leave size unchanged.

Quit is a function that prints "quitting" (in debug mode), and calls `exit(0)` to quit the program.

The rest of the functions will be written in the next tasks. The menu should be extensible, you will change and extend it in each sub-task of task 1. It should be printed using a loop iterating over the menu array, and be {NULL, NULL} terminated.

Be sure to implement this code and test it carefully before the lab (that is why you have the debug option), as you will need to extend it during the lab!

Task 1: hexeditplus

In this task we will write our own version of *hexedit* for working with binary files. You will extend your code from task 0b.
Note: You should verify that there is no error when opening a file. In case of an error, you should print a message and abort the rest of the operation.
For this task you will be working with the following ELF file: [abc](abc).

Note: For any funcitons that handle files, the file needs to be opend and closed **withing that function**

Task 1a: Load Into Memory

For this task, you may assume that the size of the file is less than 10k bytes.

Write the function for the "Load Into Memory" option, which works as follows:

- Check if *file_name* is empty (i.e. equals to ""), and if it is print an error message and return.
- Open *file_name* for reading. If this fails, print an error message and return.
- Prompt the user for *location* (in hexadecimal) and *length* (in decimal).
- If debug flag is on, print the file_name, as well as *location*, and *length*.
- Copy *length * unit_size* **bytes** from *file_name* starting at position *location* into *mem_buf*.
- Close the file.

For example, the command:

```
./hexeditplus
```
Will print:

```
Choose action:
0-Toggle Debug Mode
1-Set File Name
2-Set Unit Size
3-Load Into Memory
4-Quit
```

Assume that the user has already set the file name to "abc". If the user chooses 4, he is prompted for *location* and *length*. It should look as follows:

```
3
Please enter <location> <length>
12F 10
```

The program should open the file abc and load the 10 bytes (Assuming unit size is set to 1), from byte 303 to byte 313 in the file into *mem_buf*. The output should look like:

```
Loaded 10 units into memory
```

> **Remember**
> - To read *location* and *length* use `fgets` and then `sscanf`, rather than `scanf` directly.
> - *location* is entered in hexadecimal representation.

> **Working with units**
> You are required to write code that handles data in unit sizes (i.e. not necessarily single bytes). This might confuse you into writing much more code than needed. See this example on how to handle multiple unit sizes when reading, printing and writing without writing too much code. Relevant to this task is the function `read_units_to_memory`

Task 1b: Toggle Display Mode

Write the function for the "Toggle Display Mode" option, which swiches between display using a decimal representation, and by using a hexadecimal representation.

Toggle display mode means turn the display flag on and print using a hexadecimal representation, the initial state is off and print using a decimal representation. Print "Display flag now on, hexadecimal representation", and if the display flag is on, this function prints "Display flag now off, decimal representation", and turns the flag off. For example, the command:

```
./hexeditplus
```

Will print:

```
Choose action:
0-Toggle Debug Mode
1-Set File Name
2-Set Unit Size
3-Load Into Memory
4-Toggle Display Mode
5-Quit
```

Assume that the display flag is off.

```
4
Display flag now on, hexadecimal representation
4
Display flag now off, decimal representation
```

Task 1c: Memory Display

Write the function for the "Memory Display" option:
This option displays *u* units of size *unit_size* starting at address *addr* in memory. Unit_size is already defined in *state*, but *u* and *addr* should be queried from the user by this function. *u* will be given in decimal and *addr* in hexadecimal. Entering a value of 0 for *addr* is a special case, in which the memory to be displayed starts at your mem_buf.

The units should be displayed according to the display flag. If the display flag is on then print using a hexadecimal representation, and if the display flag is off print using a decimal representation.

For example, the command line:

```
 ./hexeditplus
```

Will print:

```
Choose action:
0-Toggle Debug Mode
1-Set File Name
2-Set Unit Size
3-Load Into Memory
4-Toggle Display Mode
5-Memory Display
6-Quit
```

If the user set the unit size to 2 and loaded a file into memory, then the output should look something like this (for a unit size of 2):

```
Choose action:
0-Toggle Debug Mode
1-Set File Name
2-Set Unit Size
3-Load Into Memory
```

```
4-Toggle Display Mode
5-Memory Display
6-Quit
> 5
Enter address and length
> 0 5
Decimal
=======
256
0
12032
54528
37120
Choose action:
0-Toggle Debug Mode
1-Set File Name
2-Set Unit Size
3-Load Into Memory
4-Toggle Display Mode
5-Memory Display
6-Quit
> 4
Display flag now on, hexadecimal representation
Choose action:
0-Toggle Debug Mode
1-Set File Name
2-Set Unit Size
3-Load Into Memory
4-Toggle Display Mode
5-Memory Display
6-Quit
> 5
Enter address and length
> 0 5
Hexadecimal
===========
100
0
2F00
D500
9100
```

- Note that, depending on the chosen unit size, the printed hexadecimal values may differ in order when compared with the output of *hexedit*. Why is that?

Use your newly implemented functionality to answer: what is the entry point of your own *hexeditplus* program? Verify your answer using `readelf -h`

> **Working with units**
> You are required to write code that handles data in unit sizes (i.e. not necessarily single bytes). This might confuse you into writing much more code than needed. See [this example](#) on how to handle multiple unit sizes when reading, printing and writing without writing too much code. Relevant to this task is the function `print_units`

Task 1d: Save Into File

Write the function for the "Save Into File" option, which works as follows:
This option replaces *length* units at *target-location* of *file_name* with bytes from the **hexeditplus** memory starting at virtual address *source-address*.

For example, the command:

```
  ./hexeditplus
```

Will print:

```
Choose action:
0-Toggle Debug Mode
1-Set File Name
2-Set Unit Size
3-Load Into Memory
4-Toggle Display Mode
5-Memory Display
6-Save Into File
7-Quit
```

When the user chooses option 6, the program should query the user for:
- *source-address* (source memory address, in hexadecimal), *source-address* can be set to `0`, in which case, the source address is start of *mem_buf*, on any other case, take an address in memory.
- *target-location* (target file offset, in hexadecimal),
- *length* (number of units, in decimal).

Implement the checks that the file can be opened (for writing and NOT truncating), and print appropriate debug messages in debug mode as in the previous task. Close the file after writing.

For example, after the file name was set to "abc" and unit size to 1 bytes, choosing option "5-Save Into File" using *source-address* 960c170, *target-location* 33 and *length* 4, the program should write *length* = 4 bytes from memory, starting at address 0x960c170 to the file *abc*, starting from offset 0x33 (overwriting what was originally there). It should look as follows:

```
6
Please enter <source-address> <target-location> <length>
960c170 33 4
```

Note that the target file is the one specified using option 1 in the menu.

Also observe that after you execute this option, **only** *length* units of the file *file_name* should be changed.

> If <target-location> is greater than the size of <file_name> you should print an error message and not copy anything.

You should use *hexedit*, to verify that your code for tasks 1c and 1d works correctly, by loading a portion of a file into memory and saving it to another file.
Here is some of *hexedit*'s output for the file abc, verify that you understand why the output is as it is.

```
00000070    01 00 00 00    01 00 00 00    00 00 00 00    00 80 04 08    ................
00000080    00 80 04 08    EC 05 00 00    EC 05 00 00    05 00 00 00    ................
00000090    00 10 00 00    01 00 00 00    14 0F 00 00    14 9F 04 08    ................
000000A0    14 9F 04 08    0C 01 00 00    14 01 00 00    06 00 00 00    ................
000000B0    00 10 00 00    02 00 00 00    28 0F 00 00    28 9F 04 08    ........(...(...
000000C0    28 9F 04 08    C8 00 00 00    C8 00 00 00    06 00 00 00    (...............
000000D0    04 00 00 00    04 00 00 00    48 01 00 00    48 81 04 08    ........H...H...
000000E0    48 81 04 08    44 00 00 00    44 00 00 00    04 00 00 00    H...D...D.......
000000F0    04 00 00 00    51 E5 74 64    00 00 00 00    00 00 00 00    ....Q.td........
00000100    00 00 00 00    00 00 00 00    00 00 00 00    06 00 00 00    ................
00000110    04 00 00 00    52 E5 74 64    14 0F 00 00    14 9F 04 08    ....R.td........
00000120    14 9F 04 08    EC 00 00 00    EC 00 00 00    04 00 00 00    ................
00000130    01 00 00 00    2F 6C 69 62    2F 6C 64 2D    6C 69 6E 75    ..../lib/ld-linu
00000140    78 2E 73 6F    2E 32 00 00    04 00 00 00    10 00 00 00    x.so.2..........
00000150    01 00 00 00    47 4E 55 00    00 00 00 00    02 00 00 00    ....GNU.........
00000160    06 00 00 00    0F 00 00 00    04 00 00 00    14 00 00 00    ................
00000170    03 00 00 00    47 4E 55 00    C1 4E 4D 18    B9 A6 21 8F    ....GNU..NM...!.
```

> **Working with units**
> You are required to write code that handles data in unit sizes (i.e. not necessarily single bytes). This might confuse you into writing much more code than needed. See [this example](#) on how to handle multiple unit sizes when reading, printing and writing without writing too much code. Relevant to this task is the function `write_units`

Task 1e: Memory Modify

Write the function for the "Memory Modify" option:
This option replaces a unit at *location in memory buffer with val.*
The steps are:

1. *Prompt the user for location and val (all in hexadecimal).*

2. *If debug mode is on, print the location and val given by the user.*

3. *Replace a unit at location in the memory with the value given by val.*

For example, the command:

```
./hexeditplus
```

Will print:

```
Choose action:
0-Toggle Debug Mode
1-Set File Name
2-Set Unit Size
3-Load Into Memory
4-Toggle Display Mode
5-Memory Display
6-Save Into File
7-Memory Modify
8-Quit
```

When the user chooses option 7, the program should query the user for:
- *location (memory buffer location, in hexadecimal)*
- *val (new value, in hexadecimal)*

For example, if unit size was set to 4, choosing option "7-Memory Modify" using location 0x40, val 0x804808a, will overwrite the 4 bytes starting at location 0x40, with the new value 804808a. It should look as follows:

```
7
Please enter <location> <val>
40 804808a
```

As in the previous task, you should check that the location chosen to be modified, given the currentunit size, is valid, and act accordingly.

You can test the correctness of your code using hexedit

Task 2: Reading ELF

Task 2a

Download the following file: deep_thought.

`deep_thought` *is an executable ELF file. It does not run as expected. Your task is to understand the reason for that.*

Do the following:

1. *Run the file.*
2. *Which function precedes main in execution ? (hint: The assembly code in Lab 4   ).*
3. *What is the virtual address to which this function is loaded (hint: use* `readelf -s`)

Task 2b

Use your hexeditplus program from task 1 to display the entry point of a file.

What are the values of location/length? How do you know that?

Use the edit functions from hexeditplus program to fix the deep_thought file, so that it behaves as expected.

Task 3: Delving Deeper into the ELF Structure

task 3a

The goal of this task is to display the compiled code (in bytes) of the function `main`, *in the abc executable above.*

In order to do that, you need to:

1. *find the offset (file location) of the function* `main`.
2. *find the size of the function* `main`.
3. *use your hexeditplus program to display the content of that function on the screen.*

Finding the needed information:

1. *Find the entry for the function* `main` *in the symbol table of the ELF executable (* `readelf -s` *).*
2. *In that reference you will find both the size of the function and the function's virtual address and section number.*
3. *In the section table of the executable, find the entry for the function's section (* `readelf -S` *).*
4. *Find both the section's virtual address (Addr), and the section's file offset (Off).*
5. *Use the above information to find the file offset of the function.*

Task 3b-asm

> This task is only for students of the architecture and splab course.

What are the first two machine instructions in function main, stated in assembly language? I.e. you need to manually dis-assemble these first two instructions.
You can use the opcode information in the [nasm manual](#) .

> Have a look at [Practical session 3](#)   page 9, before you delve into nasm manual.

Task 3b-splab

> This task is for students registered for splab course only and NOT architecture.

Hack this executable file so that it does nothing when it is run: replace the code of the `main` *function by [NOP](#)   instructions.*
Make sure you do NOT override the `ret` *instruction ([Opcode](#) : c3) in* `main`.
Alternately, you can plant just one `ret` *instruction (where?).*

Task 4: Hacking: installing a patch using hexeditplus

The following file [ntsc](#) was meant to be a digit counter. Download it, and run it in the command-line.
```
./ntsc aabbaba123baacca
./ntsc 1112111
```

What is the problem with the file? (hint, try this string: 0123456789)

Create a new program with a correct digit counter function (should get a char* and return an int), compile and test it. (remember to compile with the -m32 flag in order to produce an ELF compatible with 32bits). **You must compile with the flag `-fno-pie`**

Use hexeditplus to replace (patch) the buggy `digit_cnt` function in the ntsc file with the corrected version from the new program.
You should do it using options 4 & 5 in hexeditplus.
(think: are there any kinds of restrictions on the code you wrote for the `digit_cnt` function?)
Explain how you did it, and show that it works.

Deliverables:

Tasks 1,2, and 3a must be completed during the regular lab. Tasks 3b and 4 may be done in a completion lab, but only if you run out of time during the regular lab. The deliverables must be submitted until the end of the lab session.
You must submit source files for task 1 and task 4 and a makefile that compiles them. The source files must be named task1.c, task4.c, and makefile.

Submission instructions

- *Create a zip file with the relevant files (only).*
- *Upload zip file to the submission system.*
- *Download the zip file from the submission system and extract its content to an empty folder.*
- *Compile and test the code to make sure that it still works.*